






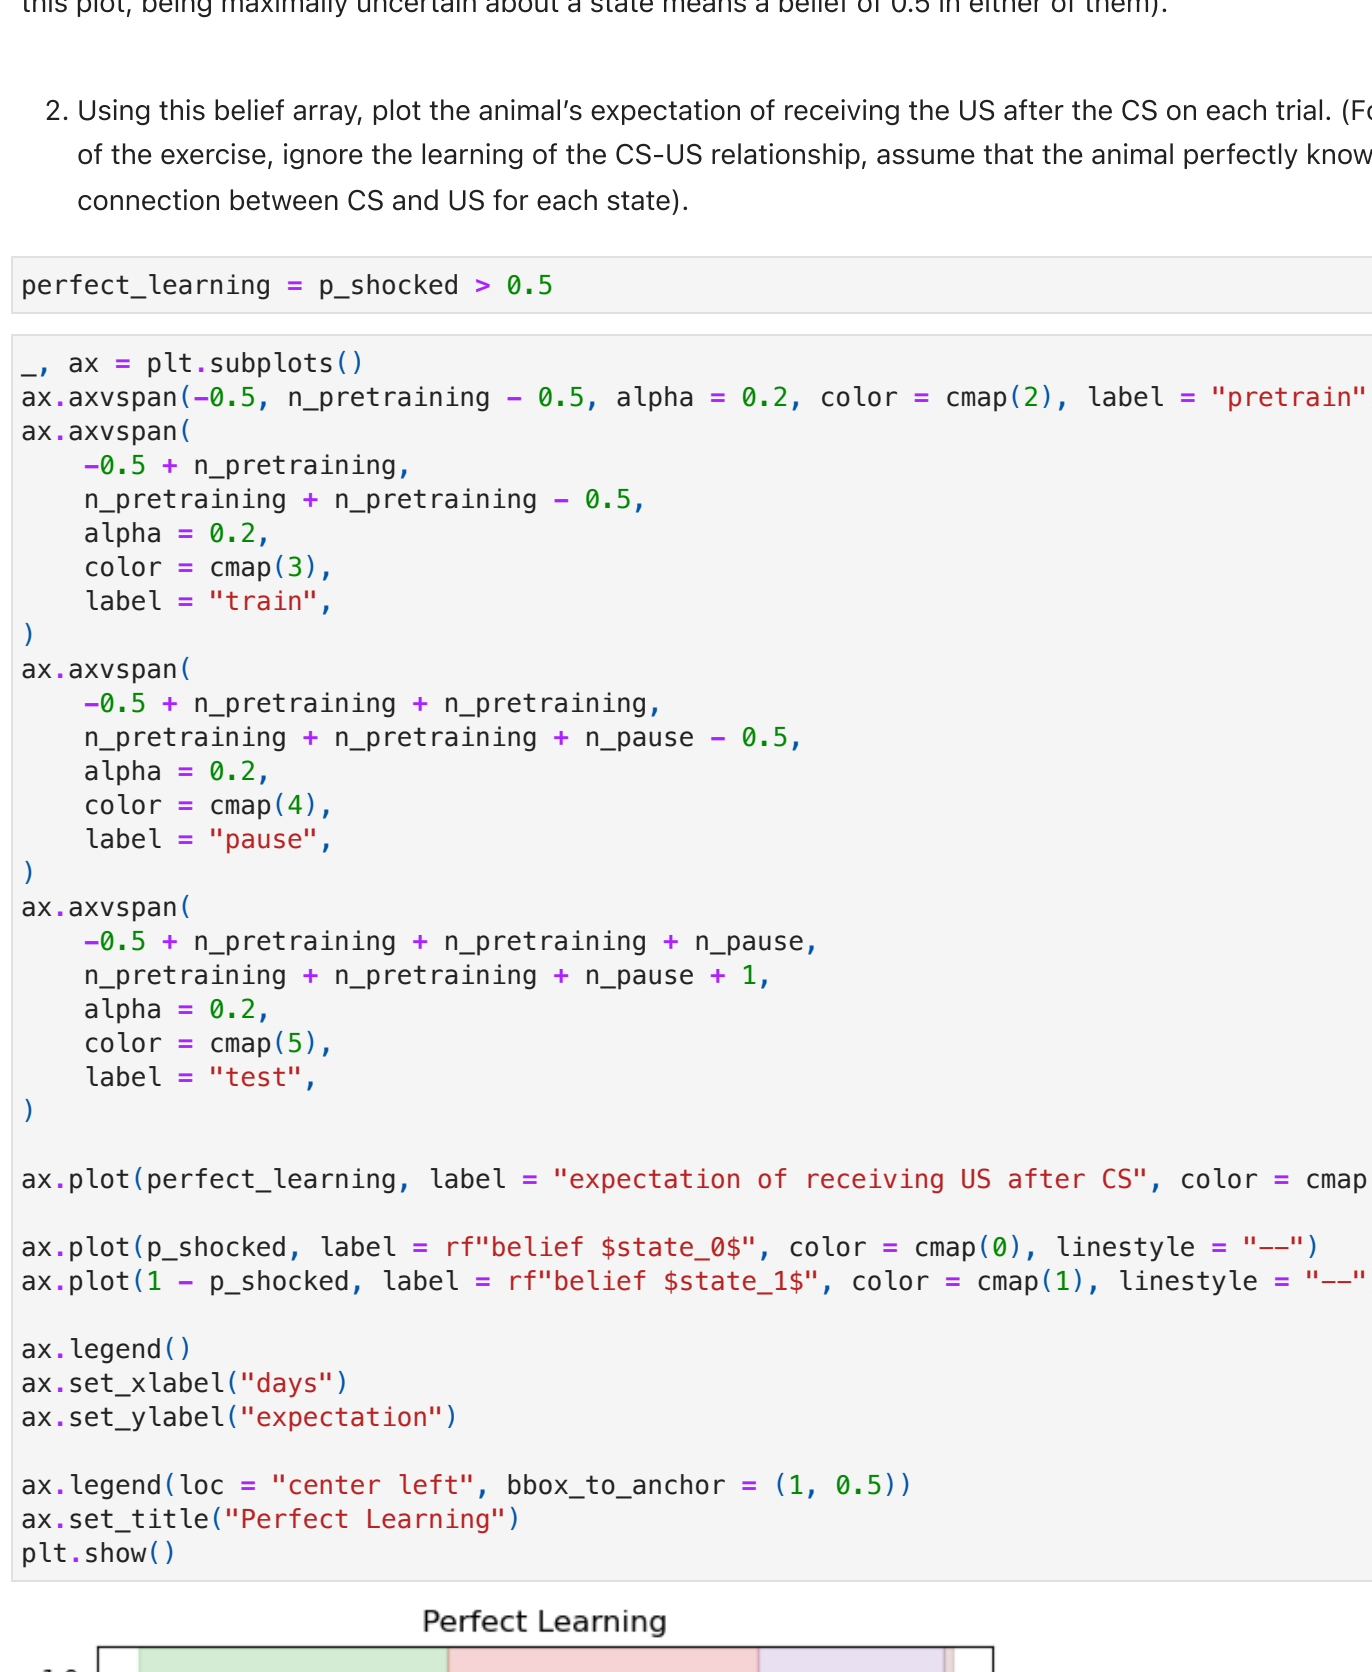
```
In [12]: pretraining = np.ones(n_pretraining)
training = np.zeros(n_training)
test = np.zeros(1)
shock = np.concatenate([pretraining, training])
tone = np.concatenate([np.ones_like(pretraining), np.zeros_like(training)], )

fig, ax = plt.subplots(figsize = (8, 5))
mapable = ax.imshow(np.stack([tone, shock]), cmap = "cividis", vmin = 0, vmax = 1)
ax.set_aspect(5)
ax.set_yticks(np.arange(0, 2), labels = ["Tone", "Shock"])
ax.set_xlabel("trials")
fig.colorbar(mapable, location = "right", label = "stimulus strength", fraction = 0.009)
ax.set_title("Stimulus")
ax.axvline(n_pretraining, color = "gray", label = "training start", linestyle = "----")
ax.legend(loc = "center left", bbox_to_anchor = (1.2, 0.5))
plt.show()
```



```
In [13]: beta = 0.7 # learning rate
p_shocked = np.zeros(n_training)
p_shocked = np.concatenate([
    np.ones(n_pretraining),
    beta * 2 * np.cos(np.arange(n_training) / n_training * np.pi) + 1 - (beta / 2),
    -(beta - 0.5) / 2 + np.cos(np.arange(n_pause + 1) / n_pause * np.pi) + (1 - beta) + (beta - 0.5)
])

fig, ax = plt.subplots(p_shocked)
ax.set_title("Expected belief")
ax.legend(loc = "center left", bbox_to_anchor = (1, 0.5))
plt.show()
```



During the pre-training phase, animal already believes that it is in the state 0. Nothing changes, until it is presented with a tone and no shock at trial 51 (start of training phase). The animal observes this change in the environment and starts to consider a possibility of a new state (state 1). As the training trials progress, its belief about being in the state 1 increases, while the belief about being in state 0 decreases. After the training, during the pause period, the animal is presented with no stimuli and the uncertainty of being in one state or the other increases (in the context of this plot, being maximally uncertain about a state means a belief of 0.5 in either of them).

2. Using this belief array, plot the animal's expectation of receiving the US after the CS on each trial. (For this part of the exercise, ignore the learning of the CS-US relationship, assume that the animal perfectly knows the connection between CS and US for each state).

```
In [14]: perfect_learning = p_shocked > 0.5

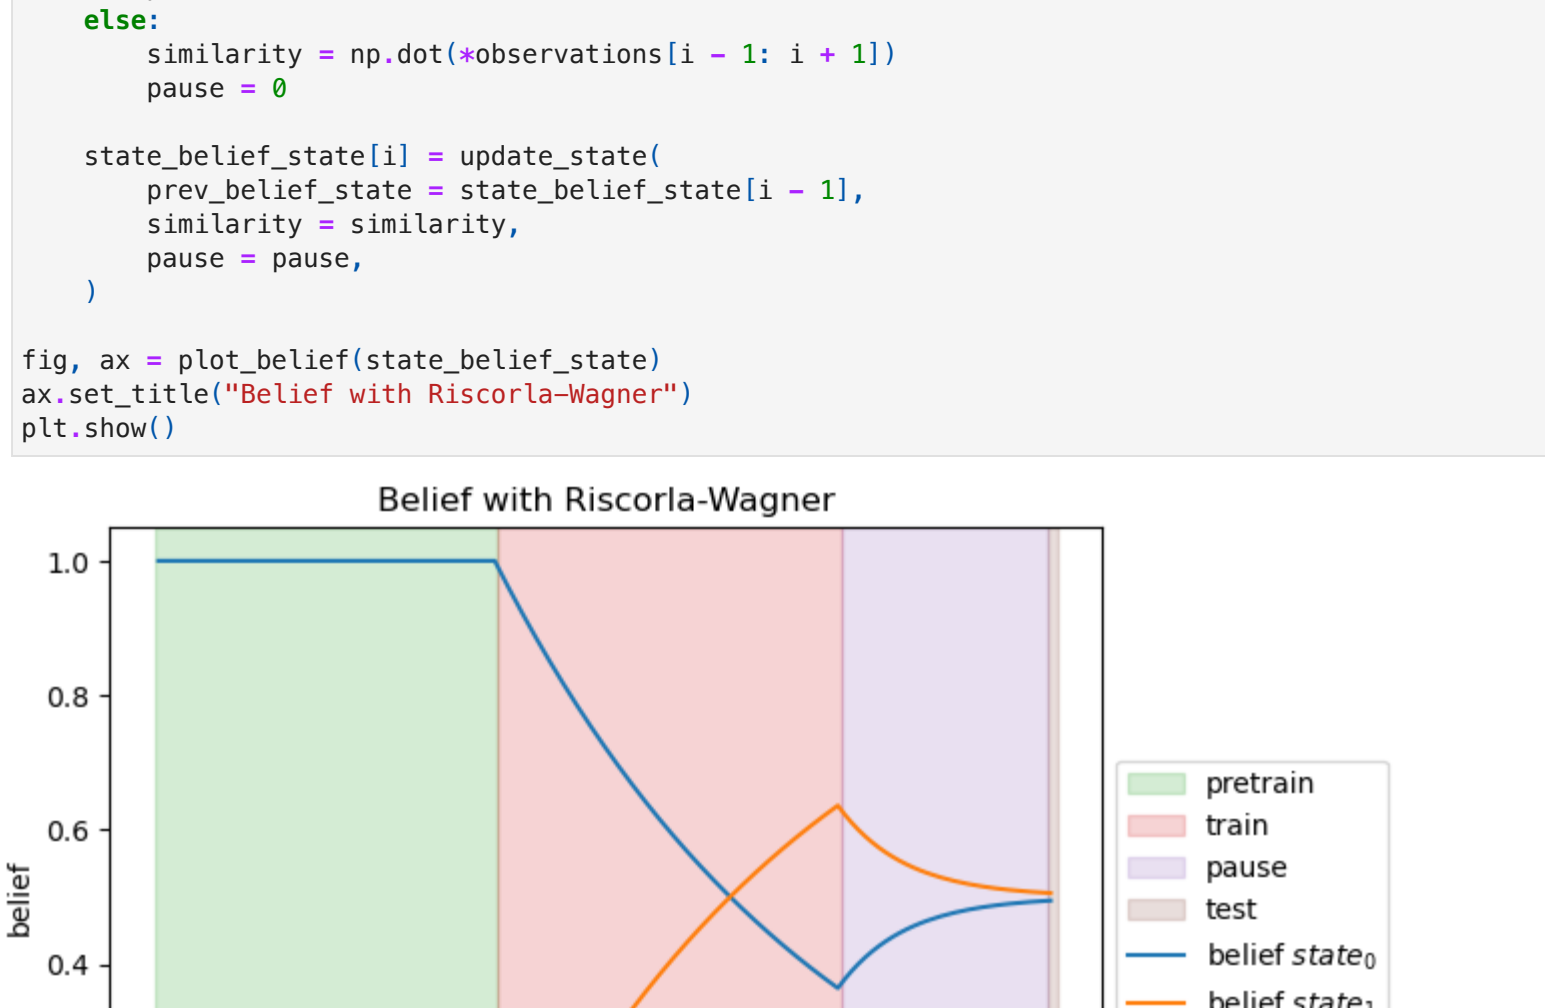
In [15]: _, ax = plt.subplots()
ax.axvspan(-0.5, n_pretraining - 0.5, alpha = 0.2, color = cmap(2), label = "pretrain")
ax.axvspan(-0.5, n_pretraining,
    alpha = 0.2,
    color = cmap(3),
    label = "train",
)
ax.axvspan(-0.5, n_pretraining + n_pretraining,
    alpha = 0.2,
    color = cmap(4),
    label = "pause",
)
ax.axvspan(-0.5, n_pretraining + n_pretraining + n_pause + 1,
    alpha = 0.2,
    color = cmap(5),
    label = "test",
)

ax.plot(perfect_learning, label = "expectation of receiving US after CS", color = cmap(2))

ax.plot(p_shocked, label = rf"belief $state_0$", color = cmap(0), linestyle = "----")
ax.plot(1 - p_shocked, label = rf"belief $state_1$", color = cmap(1), linestyle = "----")

ax.legend()
ax.set_xlabel("days")
ax.set_ylabel("expectation")

ax.legend(loc = "center left", bbox_to_anchor = (1, 0.5))
ax.set_title("Perfect Learning")
plt.show()
```



Since we assume that the animal perfectly knows the connection between CS and US for each state, the expectation of receiving the US after the CS suddenly drops to zero after the belief of being in state 1 exceeds the belief of being in state 0.

3. Now, rather than assuming our state belief array, we will infer it ourselves with a simple heuristic. Write a function which takes three inputs: The state of the previous trial, the similarity of the previous trial to the one before it in terms of made observations, and the amount of time which passed since the last trial. Based on these arguments, return a belief over all states which are under consideration. On the first trial, the animal is 100% certain of being in state 1. (We don't expect any specific functional form for this function, but it should qualitatively recreate your belief array from the first part of this exercise).

```
In [16]: def update_state(prev_belief_state: np.ndarray, similarity: float, pause: int = 0, learning_rate =
forget_rate = 0.1
for _ in range(pause):
    delta = forget_rate * (0.5 - prev_belief_state)
    prev_belief_state += delta

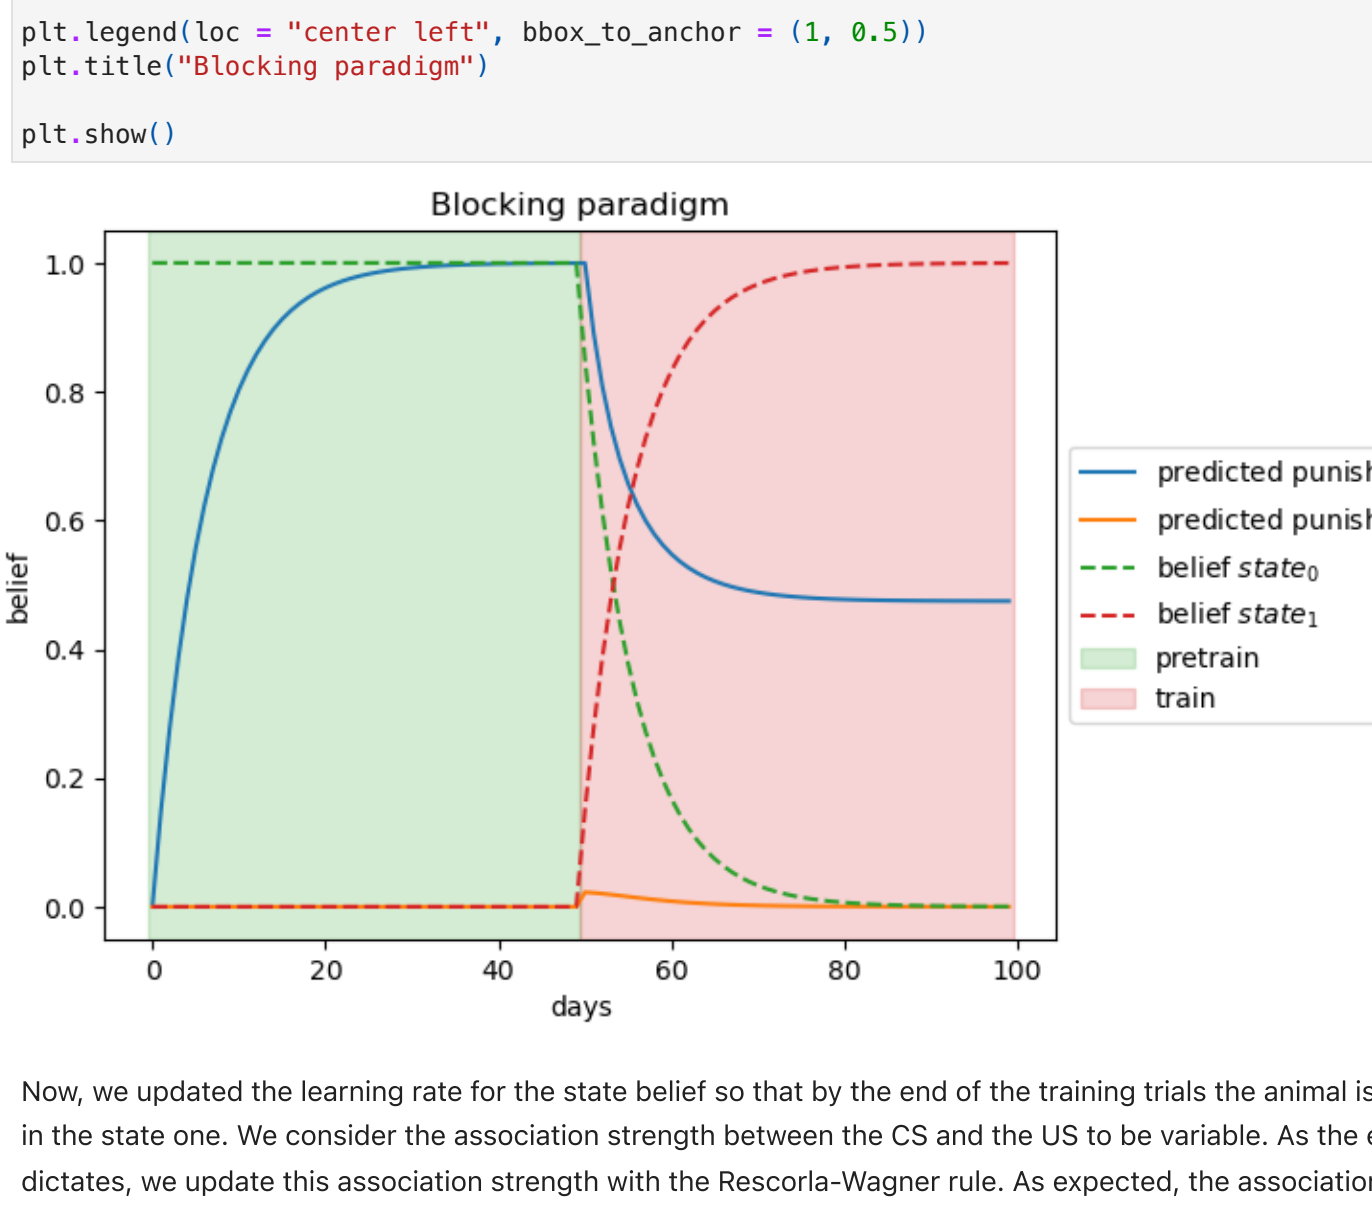
# assume 2 dimensional observation space
delta = learning_rate * (similarity - 1 - prev_belief_state)
prev_belief_state += delta

return prev_belief_state

observations = np.stack([tone, shock]).T
state_belief_state = np.empty((n_training + n_pretraining + n_pause + 1))
state_belief_state[0] = 1
for i in range(1, n_pretraining + n_training + n_pause + 1):
    # just for displaying purposes -> same thing as just putting in the 30 day pause into update_state
    if i == n_pretraining + n_training:
        similarity = state_belief_state[i - 1] * 1
        pause = 1
    else:
        similarity = np.dot(observations[i - 1: i + 1])
        pause = 0

state_belief_state[i] = update_state(
    prev_belief_state = state_belief_state[i - 1],
    similarity = similarity,
    pause = pause,
)

fig, ax = plt.subplots(state_belief_state)
ax.set_title("Belief with Rescorla-Wagner")
plt.show()
```



During the pre-training phase, the animal already believes that it is in the state 0. Nothing changes, until it is presented with CS and no US at trial 51 (start of training phase). Now, we apply the Rescorla-Wagner rule to update the belief about being in each state based on its similarity to the previous state. So, the belief about being in the state 1 increases, while the belief about being in state 0 decreases. After 50 training trials, the animal enters the pause phase, when it is not exposed to any of the stimuli. We model this using the Rescorla-Wagner rule, this time assuming that the belief about each state approaches equal probabilities (in the case of two stimuli, this probability is 0.5).

This plot looks a bit different with respect to our expectation. We used a cosine function to come up with our expectation, however, this in turn causes it to be flawed as the rate of change in belief decreases before entering a new phase (clearly visible before the pause stage).

4. Finally, maintain a learned association strength between CS and US for each state separately. Update them after each trial according to the Rescorla-Wagner rule, but also weigh the magnitude of the update by the strength of the belief in the state. For simplicity, assume that the belief about the current state does not change between the beginning and the end of each trial.

```
In [17]: # Define stimuli
tone = np.ones(n_pretraining + n_training).astype(bool)

# Define rewards
rewards = np.concatenate([
    np.ones(n_pretraining),
    np.zeros(n_training)
]).astype(bool)

In [18]: _, ax = plt.subplots(figsize = (8, 5))

plt.imshow([tone, rewards], aspect = "equal")

ax.set_aspect(5)
ax.set_yticks(np.arange(0, 2), labels = ["CS", "US"])
plt.title("Stimulus and reward presentation")
plt.show()
```



```
In [19]: ### More complicated Rescorla-Wagner (for two states, weighted by state belief)

# stimuli
u = tone
alpha = learning_rate
r = shock

# predicted rewards
vA = np.zeros(shape = (n_pretraining + n_training, 1))
vB = np.zeros(shape = (n_pretraining + n_training, 1))
# weights for state A -> tone -> shock
wA = np.zeros(shape = (n_pretraining + n_training, 1))
wB = np.zeros(shape = (n_pretraining + n_training, 1))

# Belief about the state A
state_belief_state = np.empty((n_training + n_pretraining))
state_belief_state[0] = 1

observations = np.stack([tone, shock]).T

for i in range(1, n_pretraining + n_training):
    similarity = np.dot(observations[i - 1: i + 1])
    state_belief_state[i] = update_state(
        prev_belief_state = state_belief_state[i - 1],
        similarity = similarity,
        learning_rate = 0.15
    )

    v_exp = np.sum(vA[i - 1])
    delta_w = alpha * (r[i - 1] - v_exp) * u[i - 1]
    wB[i] = wB[i - 1] + delta_w * state_belief_state[i]
    vA[i] = wA[i] * u[i]

    v_exp = np.sum(vB[i - 1])
    delta_w = alpha * (r[i - 1] - v_exp) * u[i - 1]
    wB[i] = wB[i - 1] + delta_w * (1 - state_belief_state[i])
    vB[i] = wB[i] * u[i]

In [20]: fig, ax = plt.subplots()
```



Now, we updated the learning rate for the state belief so that by the end of the training trials the animal is certain it is in the state one. We consider the association strength between the CS and the US to be variable. As the exercise dictates, we update this association strength with the Rescorla-Wagner rule. As expected, the association strength in state zero decreases because the probability of being in state zero is still quite high when the CS is presented with no US. A similar connection we can see for state one. Despite the animal believing that it is in state zero it also updates its association on not receiving a punishment in state one because there is now a chance that it might have changed the state.